

AI Go-To-Market Strategist for LinkedIn — Product Strategy & MVP Spec

1. Vision

Most people trying to build a business through social selling on LinkedIn know they *should* be posting, connecting, and building a persona — but they have no idea what to post, who to target, or how to turn attention into pipeline. This product is an AI engine that functions as a dedicated go-to-market strategist and social media strategist for a single user (or eventually a team), living on top of their LinkedIn presence.

It does not just generate a content calendar. It builds and continuously refines a *strategy* — persona, audience, positioning, content, and (eventually) outreach — based on:

1. Who the user actually is and what they're trying to achieve (captured via a structured consulting-style intake)
2. Who the user already knows (their existing LinkedIn network — someone with 100 connections needs a fundamentally different strategy than someone with 10,000)
3. What's working in their specific niche right now (industry trend and competitor/peer analysis)
4. What the user is actually doing in their day-to-day work (meeting notes, insights, documents fed back into the system)

The end state (phase 3) is a model where the engine doesn't just build a persona — it identifies a scalable business model *around* that persona, in the spirit of the commercial real estate example: someone builds enough authority and inbound interest via LinkedIn that they can recruit and fund others to replicate the model on their behalf, creating a self-reinforcing funnel of talent, deal flow, and content.

Core product promise: the user does their day job. The engine works in the background, learns continuously, proposes strategy and content, and — once trusted enough — starts running key parts of the process independently.

2. Guiding Principles

- **Strategy before content.** Nothing gets posted without an underlying “why” tied to audience and business goal.
 - **Network-aware, not generic.** Strategy must adapt to the size and shape of the user’s actual LinkedIn network, not a one-size-fits-all playbook.
 - **Human-in-the-loop by default, autonomous by earned trust.** The system should never assume it knows better than the user. It earns autonomy through a measurable confidence score built on approvals over time.
 - **Collaborative, not one-shot.** Draft content is a starting point for a conversation, not a final deliverable — the user should be able to iterate in natural language until it sounds like them.
 - **Reusable across niches.** The architecture should generalize — commercial real estate, book publishing, industrial automation sales, etc. — with the niche/persona layer as a variable, not hardcoded logic.
 - **Security first.** Any credential connection (LinkedIn, Google Drive, etc.) must be handled with real security layers — this is a trust-critical product from day one.
-

3. Phase 1 — MVP: The Strategist & Content Engine

This is what gets built first. **No autonomous outreach or connecting in Phase 1.** The MVP proves out the strategy and content loop only.

3.1 Step 1 — Secure LinkedIn Connection

- OAuth-based (or equivalent secure) LinkedIn account connection.
- Credential/token storage must follow security best practices (encrypted at rest, scoped permissions, no plaintext storage, revocable access).
- Initial pull: connection count, connection list/industry breakdown (to the extent LinkedIn’s API allows), profile data, posting history if available.
- This is a trust checkpoint for the user — worth being explicit in-product about what data is accessed and why.

3.2 Step 2 — Structured Adaptive Consulting Intake

An AI-led “consulting session” that establishes the strategic baseline. Two design constraints from Drew:

- **Preset framework** — a defined set of core question areas so the conversation doesn't wander.
- **Adapts in real time** — within that framework, the AI follows up dynamically based on what the user says, like a real consultant would.

Core question areas (framework, not rigid script):

1. **Who they are** — role, industry/niche, sub-niche, what they actually sell or offer.
2. **Where they are today** — current LinkedIn activity, existing content (if any), current network size/composition, current lead flow from social.
3. **Where they want to be** — business goals, target outcomes (leads, brand, funding, hiring, etc.), timeline.
4. **Who they're trying to reach** — ideal audience/persona, existing understanding of who they need in their network.
5. **Voice and constraints** — tone, things they will/won't say publicly, competitive sensitivities, compliance concerns if any.

Output of this step: a structured **strategic brief** — a persisted profile of the user's niche, goals, current state, and target state.

3.3 Step 3 — Deep-Dive Analysis

Combines the strategic brief with real data:

- **Network analysis** — size, composition, quality of existing connections relative to their stated target audience. Strategy branches meaningfully based on network size (e.g., ~100 connections vs. ~10,000 connections require different approaches — the former may need network-building as a precursor to content strategy; the latter may be ready for audience segmentation and targeted content immediately).
- **Industry/niche trend analysis** — what's currently resonating in their specific sub-niche (not just their broad industry).
- **Peer/competitor analysis** — identify who is doing well in adjacent or identical niches on LinkedIn, and what content patterns are working for them (formats, topics, cadence — not copying, but pattern-learning).

Output: a **roadmap** — the gap between current state and target state, translated into a phased content and positioning strategy.

3.4 Step 4 — Ongoing Content Engine

Once the roadmap exists, the engine moves into continuous operation:

- Produces **ready-to-post draft content** (not just topic ideas) aligned to the roadmap.
- **Conversational refinement** — user can respond in natural language ("I like this but change X," "this isn't how I'd phrase it," "make this shorter") and the engine revises. This is a back-and-forth loop, not a single generation.
- Content strategy continuously updates based on three live inputs (see 3.5).

3.5 Step 5 — Continuous Feedback & Confidence Loop

Three parallel signal sources feed the ongoing strategy:

1. **Industry/trend signal** — ongoing monitoring of what's working in the user's specific niche and among high-performing peers.
2. **Explicit user feedback signal** — every suggestion (content idea, angle, adjustment) is surfaced to the user as an approve/reject (yes/no) decision.
3. **Passive day-to-day signal** — ingestion of the user's own real-world work activity, primarily via meeting notes/transcripts (e.g., Google Drive integration pulling in Gemini-generated meeting notes, or user-submitted documents). This grounds the strategy in what the user is *actually* doing and learning, not just external trends.

Confidence Score mechanic:

- Every yes/no approval on a suggestion contributes to a running confidence score.
- As approval rate climbs, the system's confidence in acting independently increases.
- At a defined high-confidence threshold (e.g., ~90–95% approval rate over a meaningful sample size), the system proactively prompts the user: *"I'm ready to start running more independently — want to enable that?"*
- This confidence mechanic is the bridge into Phase 2 (autonomous outreach) — it should be built into the MVP even though autonomous *action* is not.

3.6 Phase 1 Explicit Non-Goals

- No autonomous outreach or connection requests.
 - No autonomous posting without user approval (even at high confidence — Phase 1 stays human-approved for publishing; autonomy scope is a Phase 2 decision).
 - No multi-platform support beyond LinkedIn.
-

4. Phase 2 — Autonomous Outreach & Network Growth

Once the strategy/content engine is proven and the confidence score model is validated:

- **Persona-matched prospecting** — engine identifies specific people/personas in the target audience worth connecting with, based on the roadmap.
 - **Outreach and connection requests** — initially human-approved per batch, eventually autonomous once confidence threshold is sustained specifically for outreach actions (may warrant its own separate confidence score, distinct from content approval).
 - **Conversation/DM assistance** — templated but personalized outreach messaging, with the same conversational refinement loop as content.
 - Consider whether outreach autonomy should require a *separate, higher* confidence bar than content autonomy, given the higher relationship risk of unsolicited outreach done wrong.
-

5. Phase 3 — The Funnel/Fund Model (Business Model Layer)

The long-term vision, inspired by the commercial real estate example: once a user has built real authority and inbound interest through their LinkedIn persona, the engine helps them identify a **replicable business model** built around that authority — e.g., identifying others in the audience who want to do what the user does, and structuring a way to fund/support/recruit them, creating a compounding pipeline of talent and opportunity that flows back to the original user.

This phase is more business-model consulting than product feature — the engine would need to:

- Identify patterns in the audience of people who are “customers” vs. people who are “aspiring operators” wanting to replicate the model.
 - Surface this distinction to the user as a strategic opportunity.
 - This is intentionally vague at this stage and should be treated as a direction, not a spec, until Phase 1 and 2 are validated.
-

6. Key Open Questions / Gaps to Resolve Before Building

- **LinkedIn API/ToS constraints** — LinkedIn’s official API is restrictive, especially around automation, outreach, and scraping connection data. Any outreach or auto-connect

functionality (Phase 2) needs a real compliance review before building — this could materially change what's feasible.

- **Google Drive / meeting notes ingestion** — what's the actual data contract? Full transcripts, summaries only, user-tagged excerpts? Privacy handling for potentially sensitive business conversations needs definition.
 - **Confidence score mechanics** — is it a single global score, or per-category (content topics vs. tone vs. timing vs., later, outreach)? Recommend per-category from the start to avoid one bad category dragging down overall trust.
 - **Multi-tenant vs. single-user architecture** — is this being built for Drew's own use first, or as a multi-user SaaS product from day one? This changes a lot about data isolation and auth design.
 - **Voice modeling** — how does the system learn and retain the user's specific voice over time, beyond in-session edits? (e.g., does it build a persistent style profile from accepted/rejected edits?)
 - **Definition of "resonates well"** — what metrics define successful content (engagement, profile views, inbound DMs, actual leads)? Needs a clear success metric to actually train the trend-analysis and confidence loop against.
-

7. Suggested Build Order for MVP (Phase 1)

1. LinkedIn secure connection + basic data pull
2. Consulting intake flow (framework + adaptive follow-up logic)
3. Strategic brief generation (structured output from intake)
4. Network + trend analysis module → roadmap generation
5. Content draft generation tied to roadmap
6. Conversational refinement loop on drafts
7. Approve/reject feedback capture → confidence score (content category only, to start)
8. Google Drive / meeting notes ingestion as a secondary signal source
9. Confidence-threshold prompt ("ready to run independently") as a stub for Phase 2 handoff